

Input/Output

Week 11 - Lecture

Devices

Team

- Instructor
 - Giancarlo Succi
- Teaching Assistants
 - Nikita Lozhnikov (also Tutorial Instructor)
 - Manuel Rodriguez
 - Shokhista Ergasheva

Sources

- These slides have been adapted from the original slides of the adopted book:
 - Tanenbaum & Bos, Modern Operating Systems: 4th edition, 2013
Prentice-Hall, Inc.and customised for the needs of this course.
- Additional input for the slides are detailed later

Introduction to I/O (1)

- In addition to providing abstractions such as processes, address spaces, and files, an operating system also controls all the computer's I/O (Input/Output) devices
- Operating system issue commands to the devices, catch interrupts, and handle errors

Introduction to I/O (2)

- OS should also provide an interface between the devices and the rest of the system that is simple and easy to use
- To the extent possible, the interface should be the same for all devices (device independence)
- The I/O code represents a significant fraction of the total operating system

I/O Devices (1)

- I/O devices can be roughly divided into two categories:
 - Block devices
 - Character devices

I/O Devices (2)

- **Block devices:**

- Store information in fixed-size blocks each one with its own address
- Common block sizes range from 512 to 65,536 bytes
- Transfers are in units of entire blocks
- It is possible to read or write each block independently of all the other ones
- Example: Hard disks, Blu-ray discs, etc.

I/O Devices (3)

- **Character devices:**
 - Deliver or accept stream of characters, without regard to block structure
 - Not addressable, does not have any seek operation
 - Example: Printers, mice, etc.

I/O Devices (4)

- Some devices do not fit in any category
- Examples:
 - Clocks are not block addressable. Nor do they generate or accept character streams
 - Memory-mapped screens and touch screens do not fit the model well either

I/O Devices (5)

Device	Data rate
Keyboard	10 bytes/sec
Mouse	100 bytes/sec
56K modem	7 KB/sec
Scanner at 300 dpi	1 MB/sec
Digital camcorder	3.5 MB/sec
4x Blu-ray disc	18 MB/sec
802.11n Wireless	37.5 MB/sec
USB 2.0	60 MB/sec
FireWire 800	100 MB/sec
Gigabit Ethernet	125 MB/sec
SATA 3 disk drive	600 MB/sec
USB 3.0	625 MB/sec
SCSI Ultra 5 bus	640 MB/sec
Single-lane PCIe 3.0 bus	985 MB/sec
Thunderbolt 2 bus	2.5 GB/sec
SONET OC-768 network	5 GB/sec

Figure 5-1. Some typical device, network, and bus data rates

Device Controllers (1)

- I/O units often consist of a mechanical component and an electronic component
- The electronic component is called the device controller or adapter. On personal computers, it is often a chip on the parent board or a printed circuit card that can be inserted into a PCIe expansion slot
- The mechanical component is the device itself

Device Controllers (2)

- The interface between the controller and the device is often a very low-level one
- A disk might be formatted with 2,000,000 sectors of 512 bytes per track. What actually comes off the drive, however, is a serial bit stream, starting with a preamble, then the 4096 bits in a sector, and finally a checksum, or ECC (Error-Correcting Code)

Device Controllers (3)

- The preamble is written when the disk is formatted and contains the cylinder and sector number, the sector size, and similar data, as well as synchronization information
- The controller's job is to convert the serial bit stream into a block of bytes and perform any error correction necessary
- The block of bytes is typically first assembled, bit by bit, in a buffer inside the controller
- After its checksum has been verified and the block has been declared to be error free, it can then be copied to main memory

Memory-Mapped I/O (1)

- Each controller has few registers that are used for communicating with the CPU
- By writing into these registers, the operating system can command the device to deliver data, accept data, switch itself on or off, or otherwise perform some action
- By reading from these registers, the operating system can learn what the device's state is, whether it is prepared to accept a new command

Memory-Mapped I/O (2)

- In addition to the control registers, many devices have a data buffer that the operating system can read and write
- For example: a common way for computers to display pixels on the screen is to have a video RAM, which is basically just a data buffer, available for programs or the operating system to write into

Memory-Mapped I/O (3)

- The issue thus arises of how the CPU communicates with the control registers and also with the device data buffers
- Two alternatives exist

Memory-Mapped I/O #1 (1)

- The first approach:
 - Each control register is assigned an I/O port number, an 8- or 16-bit integer
 - The set of all the I/O ports form the I/O port space, which is protected so that ordinary user programs cannot access it (only the operating system can)

Memory-Mapped I/O #1 (2)

- Using a special I/O instruction, the CPU can read in control register PORT and store the result in CPU register REG:
IN REG, PORT
- The CPU can write the contents of REG to a control register:
OUT PORT, REG
- In this scheme, the address spaces for memory and I/O are different, as shown in Fig. 5-2(a)

Memory-Mapped I/O #1 (3)




Figure 5-2. (a) Separate I/O and memory space

Memory-Mapped I/O #1 (4)

- The following instruction reads the contents of I/O port 4 and puts it in R0:
IN R0 , 4
- The following instruction reads the contents of R0 and puts it in memory word 4:
MOV R0 , 4

Memory-Mapped I/O #2 (1)

- The second approach:
 - This approach maps all the control registers into the memory space, as shown in Fig. 5-2(b). Each control register is assigned a unique memory address to which no memory is assigned. This system is called memory-mapped I/O
 - A hybrid scheme, with memory-mapped I/O data buffers and separate I/O ports for the control registers, is shown in Fig. 5-2(c)


Memory-Mapped I/O #2 (2)

One address space



(b)

Two address spaces



(c)

Figure 5-2. (b) Memory-mapped I/O.
(c) Hybrid.

Memory-Mapped I/O

How do these schemes actually work (1)

- In all cases, when the CPU wants to read a word, either from memory or from an I/O port, it puts the address it needs on the bus' address lines and then asserts a READ signal on a bus' control line
- A second signal line is used to tell whether I/O space or memory space is needed
- If it is memory space, the memory responds to the request

Memory-Mapped I/O

How do these schemes actually work (2)

- If it is I/O space, the I/O device responds to the request
- If there is only memory space (as in Fig. 5-2b), every memory module and every I/O device compares the address lines to the range of addresses that it services
- If the address falls in its range, it responds to the request
- Since no address is ever assigned to both memory and an I/O device, there is no ambiguity and no conflict

Memory-Mapped I/O

Advantages (1)

- If special I/O instructions are needed to read and write the device control registers, access to them requires the use of assembly code since there is no way to execute an IN or OUT instruction in C or C++. Calling such a procedure adds overhead to controlling I/O
- With memory-mapped I/O, device control registers are just variables in memory and can be addressed in C the same way as any other variables

Memory-Mapped I/O

Advantages (2)

- With memory-mapped I/O, an I/O device driver can be written entirely in C. Without memory-mapped I/O, some assembly code is needed
- With memory-mapped I/O, no special protection mechanism is needed to keep user processes from performing I/O. All the operating system has to do is to refrain from putting that portion of the address space containing the control registers in any user's virtual address space

Memory-Mapped I/O

Advantages (3)

- If each device has its control registers on a different page of the address space, the operating system can give a user control over specific devices but not others by simply including the desired pages in its page table
- With memory-mapped I/O, every instruction that can reference memory can also reference control registers

Memory-Mapped I/O

Advantages (4)

- For example, if there is an instruction, TEST, that tests a memory word for 0, it can also be used to test a control register for 0, which might be the signal that the device is idle and can accept a new command

```
LOOP: TEST PORT 4    // check if port 4 is 0
BEQ READY            // if it is 0, go to ready
BRANCH LOOP          // otherwise, continue testing
READY:
```

Memory-Mapped I/O

Disadvantages (1)

- Consider again the assembly-code loop:

```
LOOP: TEST  PORT 4 // check if port 4 is 0
      BEQ  READY    // if it is 0, go to ready
      BRANCH LOOP    // otherwise, continue testing
READY:
```

- The first reference to PORT 4 would cause it to be cached. Subsequent references would just take the value from the cache and not even ask the device. Then when the device finally became ready, the software would have no way of finding out. Instead, the loop would go on forever

Memory-Mapped I/O

Disadvantages (2)

- To prevent this situation with memory-mapped I/O, the hardware has to be able to selectively disable caching, for example, on a per-page basis.
- This feature adds extra complexity to both the hardware and the operating system, which has to manage the selective caching

Memory-Mapped I/O

Disadvantages (3)

- Second, if there is only one address space, then all memory modules and all I/O devices must examine all memory references to see which ones to respond to. If the computer has a single bus, as in Fig. 5-3a, having everyone look at every address is straightforward
- However, the trend in modern personal computers is to have a dedicated high speed memory bus, as shown in Fig. 5-3b

Memory-Mapped I/O

Disadvantages (4)

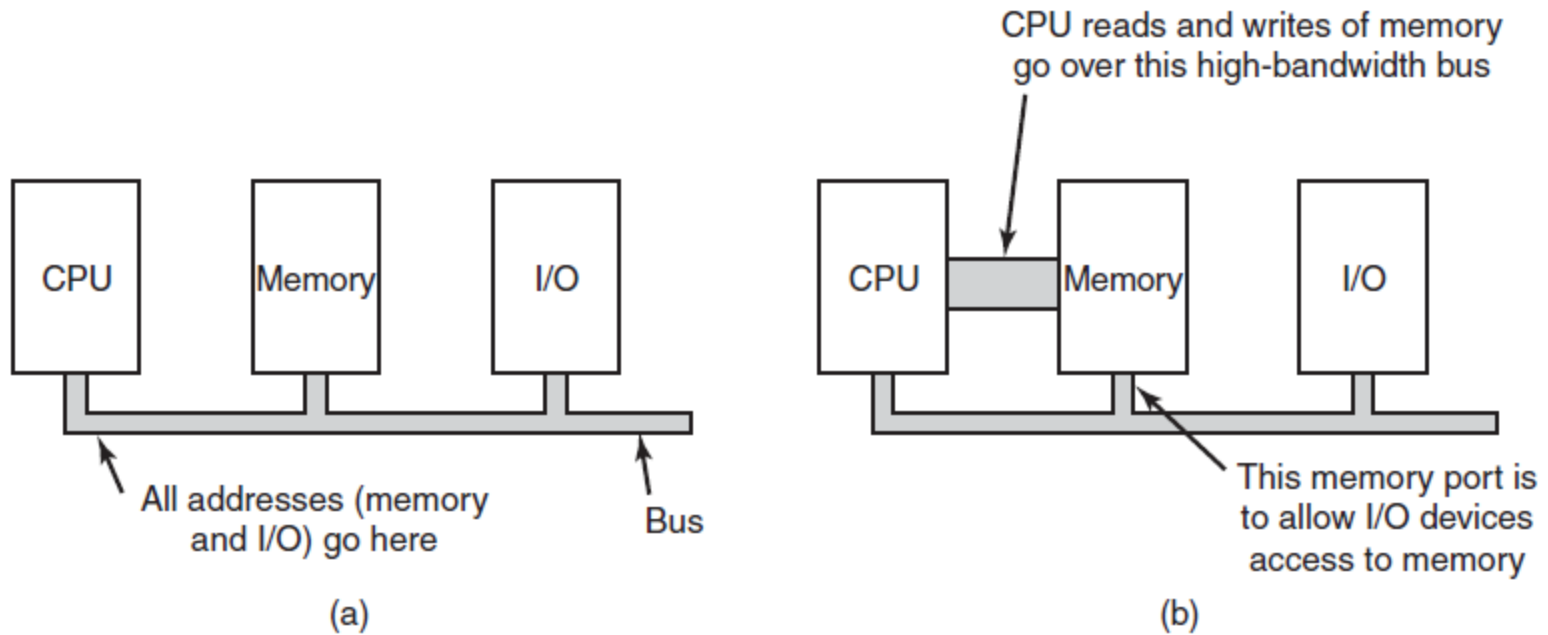


Figure 5-3. (a) A single-bus architecture
(b) A dual-bus memory architecture

Direct Memory Access (1)

- No matter whether a CPU does or does not have memory-mapped I/O, it needs to address the device controllers to exchange data with them
- The CPU can request data from an I/O controller one byte at a time, but doing so wastes the CPU's time, so a different scheme, called DMA (Direct Memory Access) is often used
- We assume that the CPU accesses all devices and memory via a single system bus that connects the CPU, the memory, and the I/O devices, as shown in Fig. 5-4

Direct Memory Access (2)




Figure 5-4. Operation of a DMA transfer

Direct Memory Access (3)

- The DMA controller has access to the system bus independent of the CPU, as shown in Fig. 5-4
- It contains several registers that can be written and read by the CPU
- These include a memory address register, a byte count register, and one or more control registers
- The control registers specify the I/O port to use, the direction of the transfer (reading from the I/O device or writing to the I/O device), the transfer unit (byte at a time or word at a time), and the number of bytes to transfer in one burst

Direct Memory Access (4)

- How Disk Read Occur when DMA is NOT used:
 - First, the disk controller reads the block (one or more sectors) from the drive serially, bit by bit, until the entire block is in the controller's internal buffer
 - Next, it computes the checksum to verify that no read errors have occurred
 - When the OS starts running, it can read the disk block from the controller's buffer a byte or a word at a time by executing a loop, with each iteration reading one byte or word from a controller device register and storing it in main memory

Direct Memory Access (5)

- How Disk Read Occur with DMA:
 - When DMA is used, the procedure is different
 - First the CPU programs the DMA controller by setting its registers so it knows what to transfer and where (Step 1 in Fig. 5-4).

Direct Memory Access (6)

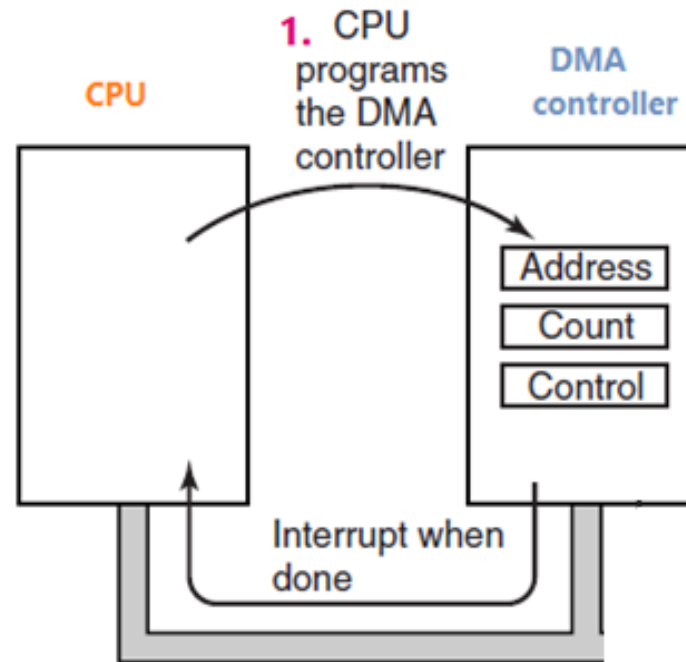


Figure 5-4. Operation of a DMA transfer. Step 1

Direct Memory Access (7)

- It also issues a command to the disk controller telling it to read data from the disk into its internal buffer and verify the checksum. When valid data are in the disk controller's buffer, DMA can begin. The DMA controller initiates the transfer by issuing a read request over the bus to the disk controller (Step 2 in Fig. 5-4)

Direct Memory Access (8)

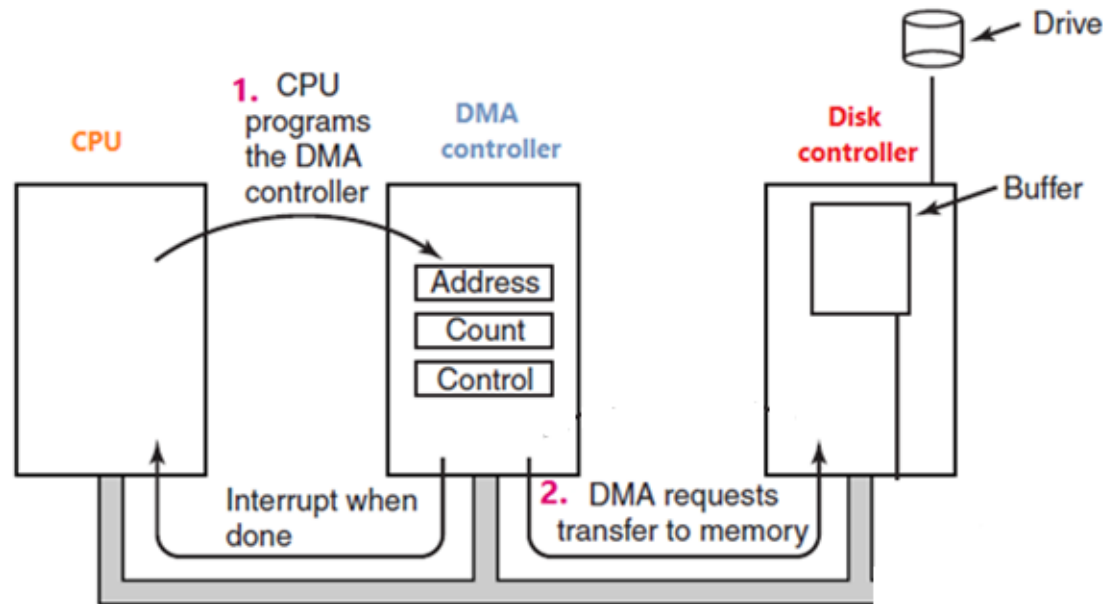


Figure 5-4. Operation of a DMA transfer. Steps 1-2

Direct Memory Access (9)

- The write to memory is another standard bus cycle (Step 3)
- When the write is complete, the disk controller sends an acknowledgement signal to the DMA controller, also over the bus (Step 4)

Direct Memory Access (10)

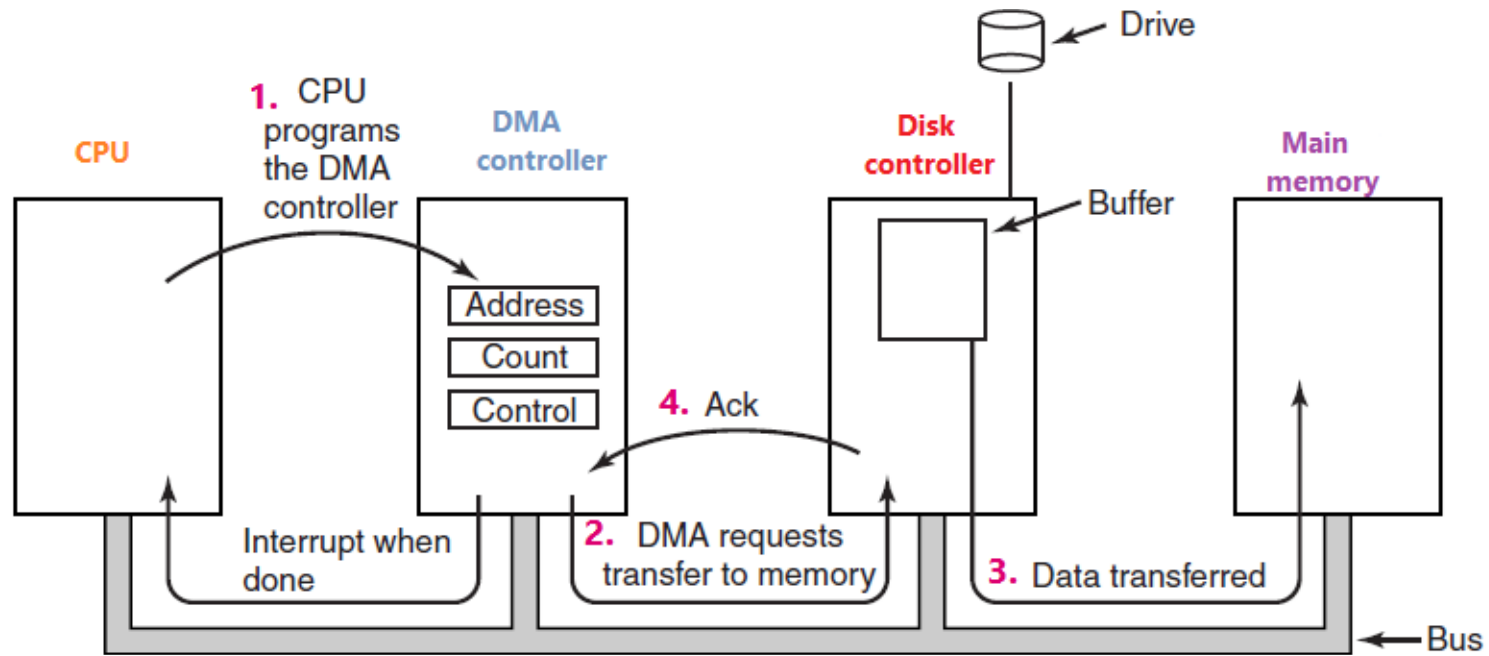


Figure 5-4. Operation of a DMA transfer. Steps 1-4

Direct Memory Access (11)

- The DMA controller then increments the memory address to use and decrements the byte count
- If the byte count is still greater than 0, steps 2 through 4 are repeated until the count reaches 0
- At that time, the DMA controller interrupts the CPU to let it know that the transfer is now complete
- When the operating system starts up, it does not have to copy the disk block to memory; it is already there

Complex DMA Controllers

- DMA controllers vary considerably in their sophistication. More complex ones can be programmed to handle multiple transfers at the same time
- Multiple requests to different device controllers may be pending at the same time, provided that there is an unambiguous way to tell the acknowledgements apart. Often a different acknowledgement line on the bus is used for each DMA channel for this reason

Bus Modes (1)

- **Word-at-a-time mode:**
 - The DMA controller requests the transfer of one word and gets it
 - If the CPU also wants the bus, it has to wait. The mechanism is called cycle stealing because the device controller sneaks in and steals an occasional bus cycle from the CPU once in a while, delaying it slightly

Bus Modes (2)

- **Block mode:**

- The DMA controller tells the device to acquire the bus, issue a series of transfers, then release the bus
- This form of operation is called burst mode
- It is more efficient than cycle stealing because acquiring the bus takes time and multiple words can be transferred for the price of one bus acquisition
- The down side to burst mode is that it can block the CPU and other devices for a substantial period if a long burst is being transferred

Bus Modes (3)

- **Fly-by mode:**

- The DMA controller tells the device controller to transfer the data directly to main memory
- An alternative mode that some DMA controllers use is to have the device controller send the word to the DMA controller, which then issues a second bus request to write the word to wherever it is supposed to go
- This scheme requires an extra bus cycle per word transferred, but is more flexible in that it can also perform device-to-device copies and even memory-to-memory copies

Transfer Using Physical Memory

- Most DMA controllers use physical memory addresses for their transfers
- Using physical addresses requires the operating system to convert the virtual address of the intended memory buffer into a physical address and write this physical address into the DMA controller's address register
- An alternative scheme used in a few DMA controllers is to write virtual addresses into the DMA controller instead.
- Then the DMA controller must use the MMU to have the virtual-to-physical translation done

END

Week 11 - Lecture

References

- Tanenbaum & Bos, Modern Operating Systems: 4th edition, 2013, Prentice-Hall, Inc.
- <http://rodrev.com/graphical/programming/Deadlock.swf>
- <http://www.utdallas.edu/~ilyen/animation/cpu/program/prog.html>
- <https://courses.cs.vt.edu/csonline/OS/Lessons/Processes/index.html>
- <http://inventwithpython.com/blog/2013/04/22/multithreaded-python-tutorial-withthreadworms/>